See launch_cpu.s and launch_gpu.s for the commands used to produce the information in C2 - C5.

C1

```
Usage:

positional arguments:
  csv_path              Path to the csv file.
  dataset_directory     Path to the dataset directory.

optional arguments:
  -h, --help            show this help message and exit
  --n_workers N_WORKERS
                        The number of workers to use for the
dataloader.
  --enable-cuda         Enable CUDA.
  --optimiser {sgd,sgdwithnesterov,adagrad,adadelta,adam}
                        The optimiser to use.
```

C2

```
Torch version 0.4.1.post2
Running on CPU.
Number of workers: 1; Optimiser: sgd
*Batch computation time includes data loading and nn forward/backward
```

Epoch	Avg Batch Load Time(s)	Avg Batch Computation Time(s)
1	3.9560	4.2445
2	3.8928	4.1808
3	3.6979	3.9969
4	3.5340	3.8283
5	3.0444	3.3335

Average time per epoch: 470.0135s

C3

Torch version 0.4.1.post2
Running on CPU.
Optimiser: sgd

Number of workers	Total Load Time (s)
0	2037.2154
1	2175.0208
2	764.8240
4	262.6946
8	33.7689
12 (best performance)	**21.3648**
16	22.6730
20	29.9539
24	29.0790
28	40.9763
32	43.9927

C4

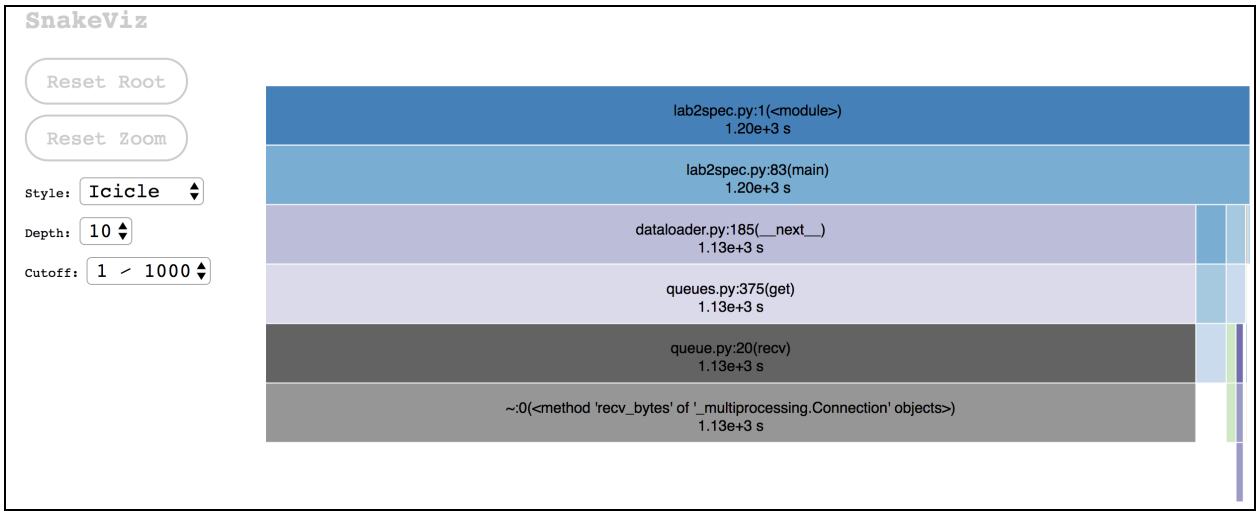

Fig. 1. Snakeviz Icicle plot, 1 worker

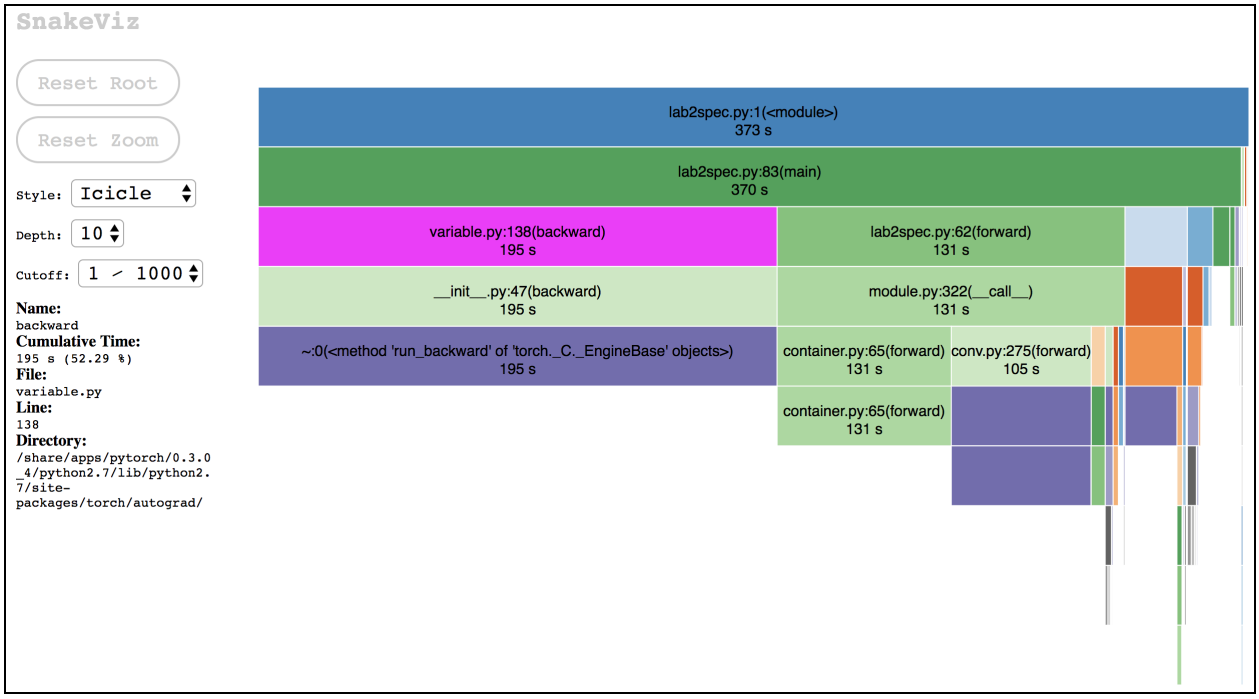

Fig. 2. Snakeviz Icicle plot, 12 workers

When using one worker, execution time is dominated by data loading. On the other hand, using 12 workers, execution time is dominated by forward/backward computations.

C5

(1)

Running on CPU.
Number of workers: 1; Optimiser: sgd
Average time per epoch: 470.0135s

Running on GPU
Number of workers: 1; Optimiser: sgd
Average time per epoch: 442.0619s

(2)

```
Torch version 0.4.1.post2
Python version 3.6.4 |Anaconda, Inc.| (default, Jan 16 2018, 18:10:19)
[GCC 7.2.0]
Running on GPU
Number of GPUs: 1
```

Number of workers: 12; **Optimiser**: sgd

Epoch	Avg Loss	Avg p@1	Avg p@3
1	0.4430	0.8324	0.5758
2	0.2758	0.9258	0.6380
3	0.2555	0.9258	0.6423
4	0.2493	0.9258	0.6433
5	0.2456	0.9258	0.6437

Average time per epoch: 150.3721s

Number of workers: 12; **Optimiser**: sgdwithnesterov

Epoch	Avg Loss	Avg p@1	Avg p@3
1	0.4030	0.8653	0.5843
2	0.2681	0.9256	0.6389
3	0.2491	0.9258	0.6417
4	0.2436	0.9258	0.6423
5	0.2398	0.9259	0.6429

Average time per epoch: 118.0847s

Number of workers: 12; **Optimiser**: adagrad

Epoch	Avg Loss	Avg p@1	Avg p@3
1	0.2493	0.9205	0.6357
2	0.2210	0.9360	0.6576
3	0.2130	0.9388	0.6663
4	0.2071	0.9398	0.6708
5	0.2033	0.9421	0.6744

Average time per epoch: 136.8236s

Number of workers: 12; **Optimiser**: adadelta

Epoch	Avg Loss	Avg p@1	Avg p@3
1	0.2765	0.9111	0.6285
2	0.2371	0.9261	0.6447
3	0.2318	0.9292	0.6480
4	0.2283	0.9326	0.6504
5	0.2245	0.9350	0.6534

Average time per epoch: 137.5252s

Number of workers: 12; Optimiser: adam

Epoch	Avg Loss	Avg p@1	Avg p@3
1	0.2714	0.9115	0.6230
2	0.2255	0.9353	0.6508
3	0.2175	0.9374	0.6579
4	0.2117	0.9409	0.6638
5	0.2063	0.9406	0.6686

Average time per epoch: 138.8749s